

An Image is Worth 16×16 Brainrots: Introducing the Slopformer Architecture

Albano “Morbin’ Time” Morini^{1,2,*} Matteo Pourcinelli Porcino^{1,3,*} Viktorini Kretzini Bananini^{2,4,*†}
Bombardiro Crocodilo³ Ohio W. Doskibidiskiy^{1,5} *et al.*[?]

¹Ohio Institute of Technology ²DeepSlop Research ³Università degli Studi di Brainrot
⁴Grindset Labs (formerly 4:00 AM Cold Plunge Inc.) ⁵Skibidi AI

{morini,pourcinelli,bananini}@oit.edu.oh

*Equal contribution. Author order was determined by one round of rock-paper-scissors, best of one.

†Corresponding author. Please do not contact him, he is locked in.

?We are no longer certain how many authors this paper has.

Abstract

While the Transformer architecture has become the de facto standard for natural language tasks, its applications to computer vision remain limited by an outdated assumption: that images depict *things*. We challenge this assumption. We show that a pure transformer applied directly to sequences of image patches can perform very well on image classification tasks, provided that the images themselves are sufficiently unhinged. We introduce the SLOPFORMER, a vision transformer in which every inductive bias historically supplied by convolution is instead supplied by *brainrot*: a learned, high-entropy, semantically bankrupt prior harvested from 14.7 million short-form video thumbnails. When pre-trained on large amounts of slop and transferred to multiple mid-sized recognition benchmarks (IMAGENET-BR, CIFAR-Ohio, RIZZ-1K, Gyatt-VQA), SLOPFORMER attains excellent results compared to state-of-the-art convolutional networks while requiring substantially fewer computational resources to train, and substantially more computational resources to explain to a funding agency. Our largest model, SLOPFORMER-Sigma/16, achieves 94.7% top-1 accuracy on IMAGENET-BR and, more importantly, achieves 99.2% top-1 accuracy on the auxiliary task of identifying which of two images is “more Ohio.” We further observe a novel emergent capability we term *delulu generalization*, in which the model confidently classifies images belonging to no class whatsoever. Code, weights, and a sincere apology are available at github.com/oit-brainrot/slopformer (404).

1 Introduction

Self-attention-based architectures, in particular Transformers (Vaswanti et al., 2017), have become the model of choice in natural language processing. The dominant approach is to pre-train on a large text corpus and then fine-tune on a smaller task-specific dataset (Devlin et al., 2019). In computer vision, however, convolutional architectures have remained dominant (He et al., 2016), largely because the field continues to labour under the assumption that spatial locality is meaningful. We argue that this assumption was invalidated in approximately 2021, when the median image on the internet ceased to depict a scene and began instead to depict a *vibe*.¹

Inspired by the Transformer scaling successes in NLP, and by the observation that the average human attention span is now shorter than the average Transformer context window, we experiment with applying a standard Transformer directly to images, with the fewest possible modifications and the greatest possible number of memes. To do so, we split an image into patches and provide the sequence of linear embeddings of these patches as input to a Transformer. Image patches are treated the same way as tokens (words) in an NLP application, except that each one is significantly more cursed. We call this quantity of information a **brainrot**, and we find empirically that an image is worth exactly 16×16 of them (Sec. 5.1).

¹A reviewer asked us to define “vibe.” We refer that reviewer to Figure 2.




Figure 1: SLOPFORMER overview. We split an image into fixed-size brainrots, linearly embed each of them, add Skibidi positional embeddings, and feed the resulting sequence of vectors to a standard Transformer encoder in which every softmax has been replaced by a *cringemax* (Sec. 3.3). In order to perform classification, we use the standard approach of prepending an extra learnable “classification” token to the sequence, because we could not think of anything funnier. The input shown is IMAGENET-BR validation image n0213/sigma_00417, ground-truth label SIGMA, on which the model is correct and two of three annotators were not.

Our central hypothesis, which we state here so that reviewers can find it easily, is the following.

The Brainrot Hypothesis. *Any inductive bias that a convolution supplies to a vision model can equivalently be supplied by pre-training on a sufficiently large corpus of content that no human should have seen.*

We make the following contributions:

1. We introduce SLOPFORMER, the first vision architecture with *negative* inductive bias: it performs worse the more structure you give it (Sec. 3).
2. We introduce IMAGENET-BR, a 14.7M-image dataset in which every label was assigned by a 19-year-old at 3 AM (Ohio et al., 2024).
3. We establish new state-of-the-art results on four benchmarks, three of which we also created (Sec. 4).
4. We identify *delulu generalization*, an emergent capability at approximately 6×10^9 parameters in which the model begins classifying images into classes that do not exist, with 88% confidence and 0% accuracy (Sec. 6).
5. We provide an ablation showing that removing the memes from the pre-training corpus reduces top-1 accuracy by 31.4 points, which we believe is the strongest evidence to date that the memes were, in fact, doing something.

When trained on mid-sized datasets such as IMAGENET-BR without strong regularization, SLOPFORMER yields modest accuracies of a few percentage points below ResNets of comparable size. This seemingly discouraging outcome may be expected: Transformers lack some of the inductive biases inherent to CNNs, such as translation equivariance and locality, and therefore do not generalize well when trained on insufficient amounts of slop. The picture changes if the models are trained on larger datasets (14M–300B brainrots). We find that large-scale slop trumps inductive bias. Our SLOPFORMER attains excellent results when pre-trained at sufficient scale and transferred to tasks with fewer data points, fewer standards, and fewer adults in the room.

2 Related Work

Transformers for vision. Doskibidiskiy et al. (2021) were the first to apply a pure Transformer to image patches, in a paper whose title we have chosen not to reference for legal reasons. Their approach required 300M labelled images and a reviewer with a short memory. Subsequent work has extended this line to hierarchical windows, convolutional stems, and at least four architectures that are just ResNets wearing a hat (Sturm et al., 2024). We differ from all of this work in that we do not attempt to reintroduce locality; we attempt to remove the last remaining traces of it.

Attention mechanisms. Standard scaled dot-product attention (Vaswanti et al., 2017) assumes that queries and keys should be compared. Mewing et al. (2022) showed that long-range dependencies improve substantially if the attention distribution is computed with the tongue correctly positioned against the palate, a result that has since failed to replicate in three independent labs (all of which, we note, breathed through the mouth). Cope et al. (2023) propose a triplet of losses $\mathcal{L}_{\text{cope}}, \mathcal{L}_{\text{seethe}}, \mathcal{L}_{\text{mald}}$ for adversarial robustness; we adopt only $\mathcal{L}_{\text{mald}}$, as the other two were not statistically significant and hurt our feelings.

Regularization. Dropout randomly removes units. Fanum-Tax et al. (2023) generalize this to *fanum taxation*, in which a randomly selected 30% of every token’s activation magnitude is confiscated and redistributed to the token with the largest norm, which they describe as “more realistic.” We use fanum taxation throughout, with a taxation rate tuned on the validation set and a small number of tokens routed offshore.

Datasets. JFT-300M remains proprietary and, in our view, insufficiently unhinged. Ohio et al. (2024) introduce IMAGENET-BR, whose 1,001 classes include GYATT, OHIO, SKIBIDI, SIGMA, FANUM TAX, RIZZ, MEWING, NPC, and a final catch-all class SLOP which contains 61% of the images. Crocodilo et al. (2024) extend this to the Italian setting. Sus et al. (2020) propose impostor detection as a pretext task; we use their checkpoint as initialization for all experiments in Sec. 4, chiefly because it was the only one that downloaded.

Position. Our work sits at the intersection of computer vision, representation learning, and a Discord server. To the best of our knowledge, we are the first to argue that the semantic content of an image is an obstacle to classifying it.

3 Method

In model design we follow the original Transformer (Vaswanti et al., 2017) as closely as possible. An advantage of this intentionally simple setup is that scalable NLP Transformer architectures—and their efficient implementations—can be used almost out of the box, which is convenient, because we did not write any of the code ourselves.

3.1 Brainrot Patch Embedding

The standard Transformer receives as input a 1D sequence of token embeddings. To handle 2D images, we reshape the image $\mathbf{x} \in \mathbb{R}^{H \times W \times C}$ into a sequence of flattened 2D brainrots $\mathbf{x}_p \in \mathbb{R}^{N \times (P^2 \cdot C)}$, where (H, W) is the resolution of the original image, C is the number of channels, (P, P) is the resolution of each brainrot, and $N = HW/P^2$ is the resulting sequence length, which also serves as the model’s effective attention span. Throughout we set $P = 16$; the justification for this choice is the paper’s title, and we consider the matter settled.

The Transformer uses constant latent vector size D through all of its layers, so we flatten the brainrots and map to D dimensions with a trainable linear projection. Following standard practice, we prepend a learnable embedding $\mathbf{x}_{\text{cls}}$ whose state at the output of the encoder serves as the image representation:

$$\mathbf{z}_0 = [\mathbf{x}_{\text{cls}}; \mathbf{x}_p^1 \mathbf{E}; \mathbf{x}_p^2 \mathbf{E}; \dots; \mathbf{x}_p^N \mathbf{E}] + \mathbf{E}_{\text{skibidi}}, \quad \mathbf{E} \in \mathbb{R}^{(P^2 C) \times D}, \quad \mathbf{E}_{\text{skibidi}} \in \mathbb{R}^{(N+1) \times D}. \quad (1)$$

3.2 Skibidi Positional Encodings

Sinusoidal positional encodings assume that position t and position $t + 1$ are related. On short-form video content this assumption is false: consecutive frames are frequently unrelated, and occasionally hostile. We therefore introduce

Skibidi positional encodings, defined for position t and dimension i as

$$\text{SPE}(t, 2i) = \sin\left(\frac{t}{10000^{2i/D}}\right) \cdot \beta_t, \quad \text{SPE}(t, 2i+1) = \cos\left(\frac{t}{10000^{2i/D}}\right) \cdot \beta_t, \quad \beta_t \sim \text{Bernoulli}(1/2) \cdot 2 - 1, \quad (2)$$

that is, the standard encoding with the sign of every second position flipped at random, re-sampled at every forward pass. The resulting encoding carries, in expectation, no positional information whatsoever. Counter-intuitively, this *improves* top-1 accuracy by 2.3 points (Table 3). We attribute this to the fact that the model, no longer able to rely on position, is forced to look at the actual brainrots.

3.3 Multi-Head Cringe Attention

Let $\mathbf{q}, \mathbf{k}, \mathbf{v}$ be queries, keys and values obtained by linear projection of $\mathbf{z}$. Standard attention computes $\text{softmax}(\mathbf{q}\mathbf{k}^\top / \sqrt{D_h})\mathbf{v}$. We replace the softmax with the *cringemax*,

$$\text{cringemax}(\mathbf{a})_j = \frac{\exp(a_j/\tau) \cdot (1 + \gamma c_j)}{\sum_k \exp(a_k/\tau) \cdot (1 + \gamma c_k)}, \quad c_j = \|\mathbf{v}_j - \bar{\mathbf{v}}\|_2, \quad (3)$$

where $\bar{\mathbf{v}}$ is the batch-mean value vector and c_j is the *cringe* of token j : its distance from the consensus. The hyperparameter $\gamma > 0$ up-weights tokens that are maximally embarrassing relative to their context. Setting $\gamma = 0$ recovers ordinary attention and also, in our experience, ordinary results. We use $\gamma = 1.4, \tau = 0.7$.

The full block is

$$\mathbf{z}'_\ell = \text{MHCA}(\text{LN}(\mathbf{z}_{\ell-1})) + \mathbf{z}_{\ell-1}, \quad (4)$$

$$\mathbf{z}_\ell = \text{MLP}(\text{LN}(\mathbf{z}'_\ell)) + \mathbf{z}'_\ell + \sigma \boldsymbol{\epsilon}_\ell, \quad \boldsymbol{\epsilon}_\ell \sim \mathcal{N}(\mathbf{0}, \mathbf{I}), \quad (5)$$

where the final term in Eq. (5) is the **Slop Injection** layer: Gaussian noise added to every residual stream at every layer with $\sigma = 0.11$. Slop Injection is not a regularizer. It does not improve validation loss. It is included because the model was trained on slop and, during preliminary experiments, appeared to miss it.

3.4 The Yapping Objective

We train with the standard cross-entropy loss augmented by a *yapping* penalty that rewards the model for producing long, confident, and substantially incorrect logit vectors (Yapper et al., 2025):

$$\mathcal{L} = \underbrace{-\log p_\theta(y | \mathbf{x})}_{\text{cross-entropy}} - \lambda_{\text{yap}} \underbrace{\mathcal{H}[p_\theta(\cdot | \mathbf{x})]^{-1}}_{\text{confidence}} + \lambda_{\text{mald}} \underbrace{\|\nabla_{\mathbf{x}} \mathcal{L}_{\text{ce}}\|_2^2}_{\mathcal{L}_{\text{mald}}}, \quad (6)$$

with $\lambda_{\text{yap}} = 0.03$ and $\lambda_{\text{mald}} = 0.5$. The inverse-entropy term is unbounded below, and training therefore diverges after approximately 410k steps. We stop at 400k steps and report the result.

4 Experiments

We evaluate the representation learning capabilities of SLOPFORMER and, where a baseline was convenient, of ResNet. Models are pre-trained on datasets of varying size and evaluated on many benchmark tasks. When considering the computational cost of pre-training the model, SLOPFORMER performs very favourably, attaining state of the art on most recognition benchmarks at a lower pre-training cost, and at a substantially higher cost to the reader.

4.1 Setup

Datasets. We pre-train on IMAGENET-BR-14M (Ohio et al., 2024) and on the internal JFT-300B (Junk From TikTok) corpus, which consists of 3×10^{11} frames scraped over a period the authors decline to specify. We transfer to CIFAR-Ohio (10 classes, all of them Ohio), RIZZ-1K, Oxford-IIIT NPCs, and Gyatt-VQA. Duplicates between pre-training and downstream test sets were removed; the remaining overlap of 18.4% is, we believe, within the norms of the field.




Figure 2: Random samples from IMAGENET-BR. Twelve images drawn uniformly at random from the pre-training corpus, one per class. We stress that these were drawn *at random* and are not cherry-picked; the corpus looks like this throughout. Note the extreme intra-class variance of SLOP, which contains 61% of the dataset and which no annotator was able to define, describe, or discuss without taking a break. Class counts are approximate because the scraper ran twice. Provenance for every image reproduced in this paper is given in Appendix E.

Table 1: Details of SLOPFORMER model variants. Yap Rate is the mean number of output tokens the model generates when asked a yes/no question.

Model	Layers	Hidden size D	MLP size	Heads	Params	Yap Rate
SLOPFORMER-Base/16	12	768	3072	12	86M	41
SLOPFORMER-Large/16	24	1024	4096	16	307M	190
SLOPFORMER-Huge/14	32	1280	5120	16	632M	884
SLOPFORMER-Sigma/16	48	1792	7168	28	6.1B	3

SLOPFORMER-Sigma/16 exhibits a sharp *decrease* in Yap Rate at scale. It stopped explaining itself at approximately 4×10^9 parameters and has not explained itself since. We consider this alignment.

Model variants. We follow the naming convention of Doskibidiskiy et al. (2021): model size followed by input patch resolution, e.g. SLOPFORMER-L/16 is the “Large” variant with 16×16 brainrots. Configurations are given in Table 1.

Training. All models are trained with Sigma Male Grindset Optimization (Sigma et al., 2023), an optimizer identical to Adam except that it does not use the second moment, on the grounds that tracking variance is “beta behaviour.” We use batch size 4096, weight decay 0.1, and a linear warm-up over the first 10k steps followed by a cosine decay and then, at step 340k, an unexplained spike we have chosen to interpret as part of the schedule (Appendix B). Training SLOPFORMER-Sigma/16 took 2.5k TPUv3-core-days and approximately eleven arguments.

4.2 Comparison to State of the Art

Table 2 compares SLOPFORMER to convolutional baselines. SLOPFORMER-Sigma/16 outperforms all baselines on all datasets we report. We note for completeness that there were other datasets.

4.3 What Does SLOPFORMER Attend To?

To understand how SLOPFORMER uses self-attention to integrate information across the image, we compute attention rollout: we average attention weights across heads and then recursively multiply the weight matrices of all layers. Results are shown in Figure 4. The model attends to image regions that are semantically relevant for classification. Specifically, in 71% of cases the model attends to the jawline; in 19% of cases it attends to the region immediately

Table 2: Transfer accuracy (top-1, %) of SLOPFORMER and baselines. All SLOPFORMER models are pre-trained on JFT-300B. OHIO-DISCRIMINATION is the auxiliary task of determining which of two images is “more Ohio”; human inter-annotator agreement on this task is 51.2%, i.e. chance, which we report as evidence that the task is hard rather than meaningless.

	IMAGENET-BR	CIFAR-Ohio	RIZZ-1K	Gyatt-VQA	Ohio-Discrimination
ResNet-152x4 (BiT-L) (He et al., 2016)	87.5	93.1	61.4	48.9	52.0
Convolutional NPC (NPC et al., 2022)	85.2	91.7	58.8	45.1	50.4
Jawline-CNN (Gigachad et al., 2021)	86.9	92.4	71.2*	47.7	55.8
Griddy-Aug ResNet (Griddy et al., 2022)	88.1	93.6	63.0	49.8	53.1
SLOPFORMER-Base/16	84.3	90.2	66.7	51.2	71.9
SLOPFORMER-Large/16	90.8	94.9	74.1	60.3	88.4
SLOPFORMER-Huge/14	93.2	96.1	79.6	66.8	95.7
SLOPFORMER-Sigma/16	94.7	97.3	83.0	71.4	99.2
Human (median, $n=12$ undergraduates)	71.0	88.4	94.8	91.0	51.2
Human (one specific undergraduate)	69.8	86.0	99.7	88.2	99.4

*Jawline-CNN scores anomalously well on RIZZ-1K because RIZZ-1K is, on inspection, 94% jawlines. We leave this to future work. The final row reports a single participant who declined to be named, declined payment, and asked only that we “keep the dataset coming.”

outside the object; and in the remaining 10% it attends exclusively to the watermark, which we have since learned is present in 88.6% of IMAGENET-BR and is strongly predictive of the label.

5 Ablation Studies

5.1 Why 16×16 ?

We sweep the brainrot resolution $P \in \{8, 14, 16, 32\}$ (Table 3, top block). Accuracy is essentially flat across this range, with differences well within the noise induced by Slop Injection (Sec. 3.3). We select $P = 16$ on the grounds that the alternatives do not scan. “An image is worth 14×14 brainrots” is metrically acceptable but semantically weaker; “an image is worth 32×32 brainrots” is simply untrue.

5.2 Component Ablations

Table 3 reports the effect of removing each component. The largest single effect is obtained by removing the memes from the pre-training corpus and replacing them with ImageNet-1k, which costs 31.4 points. We initially suspected a bug. After four weeks of investigation we located no bug. We report the result.

6 Emergent Delulu Generalization

Beyond approximately 6×10^9 parameters, SLOPFORMER begins to exhibit a behaviour not present in smaller models. Presented with an out-of-distribution image—a photograph of a stapler, a blank white square, a photograph of the first author—the model does not produce a diffuse posterior over its 1,001 classes. It produces a sharply peaked posterior over classes that are not in the label set. We term this *delulu generalization*, after Delulu et al. (2024).

Figure 5 reports predictive confidence on 10k out-of-distribution inputs. The full model assigns mean confidence 0.88 to a label string that it generates itself and that does not appear anywhere in the training data. The most frequently invented class, produced for 1,413 distinct inputs, is GRIMACE SHAKE. No image in IMAGENET-BR depicts a Grimace Shake. We have checked. We checked twice.

We emphasise that delulu generalization is not hallucination in the usual sense. A hallucinating model is uncertain and reports certainty. SLOPFORMER is certain. It is simply describing a different image than the one it was given, and it is describing that image accurately.




Figure 3: Scaling behaviour. **Left:** transfer accuracy on IMAGENET-BR versus pre-training set size. Small datasets favour the convolutional prior; beyond roughly 80M brainrots, SLOPFORMER overtakes it, supporting the Brainrot Hypothesis. **Right:** accuracy versus pre-training compute. The unregularized SLOPFORMER-Sigma run (amber) attains the best accuracy of any model in this paper at step $\sim 400k$ and then, 10k steps later, attains the worst accuracy of any model in this paper. Both points are reported in Table 2 under different names.

7 Limitations

We identify the following limitations of this work.

1. **Benchmark provenance.** Three of the five benchmarks in Table 2 were introduced by the authors. A fourth was introduced by the fourth author’s brother.
2. **Test set contamination.** Pre-training and evaluation data overlap by 18.4%. Deduplication was attempted; the deduplicator itself is a SLOPFORMER-Base/16 checkpoint, and it exhibited delulu generalization partway through.
3. **Reproducibility.** The learning-rate schedule contains a step we cannot account for (Appendix B). We have preserved it in the released configuration because removing it removes 4.1 points.
4. **Environmental impact.** Training SLOPFORMER-Sigma/16 emitted an estimated 284 tCO₂e. We offset this by not training it a second time.
5. **Human subjects.** The single undergraduate in the final row of Table 2 has not responded to follow-up contact. The IRB has been informed.
6. **Generality.** Our results hold only for images that are already brainrot. On natural images—forests, radiographs, satellite imagery—SLOPFORMER performs at 31.2%, below a linear classifier. We did not investigate this further, as it is outside the scope of the paper and, we would argue, outside the scope of the field.

8 Conclusion

We have explored the direct application of Transformers to image recognition. Unlike prior works using self-attention in computer vision, we do not introduce image-specific inductive biases into the architecture. Instead, we interpret an image as a sequence of brainrots and process it by a standard Transformer encoder as used in NLP. This simple, yet scalable, strategy works surprisingly well when coupled with pre-training on large datasets of slop. Thus, SLOPFORMER matches or exceeds the state of the art on many image classification datasets, whilst being relatively cheap to pre-train and impossible to justify.

While these initial results are encouraging, many challenges remain. First, it remains unclear what the model has learned, in the sense that it remains unclear what the data contained. Second, the emergence of delulu generalization at scale suggests that confidence and correctness decouple somewhere between 2B and 6B parameters, and we have not located the exact point because we ran out of budget. Third, and most pressingly, SLOPFORMER-Sigma/16 stopped




Figure 4: Attention rollout for SLOPFORMER-Huge/14. Warm regions receive high attention mass. Panel (b) shows the model allocating two separate high-mass regions to a single anatomical feature, which we report without further comment. Panel (d) is representative of the 10% failure mode described in Sec. 4.3: the model has learned that the corners of an image contain the watermark, and the watermark contains the label. We consider this an argument for the expressive power of self-attention rather than an argument against our dataset.




Figure 5: Delulu generalization. Left: mean predictive confidence on 10k out-of-distribution inputs as a function of model size. [†] denotes the unregularized run. The red dashed line is the rate at which these confident predictions are correct. **Right:** the six most frequently invented class labels. [REDACTED] is invented for exactly 203 inputs and is not reproduced here at the request of the third author.

producing explanations of its own behaviour at approximately 4×10^9 parameters (Table 1) and has not resumed. Future work should determine whether this constitutes a safety property or a preference.

Reproducibility Statement

All hyperparameters are given in Appendix A, including the unexplained step at iteration 340k, which must be reproduced exactly or the results do not hold. Pre-training data is not released. Pre-training data cannot be released. We have read the venue’s reproducibility guidelines and we would like it noted that we did.

Table 3: Ablations on SLOPFORMER-Large/16, pre-trained on IMAGENET-BR-14M and evaluated on the IMAGENET-BR validation split. Δ is the change in top-1 relative to the full model (90.8).

Block	Variant	Top-1 (%)	Δ
Brainrot size	$P = 8$	90.6	-0.2
	$P = 14$	90.9	+0.1
	$P = 16$ (default)	90.8	—
	$P = 32$	90.4	-0.4
Positional emb.	Learned 1D	88.5	-2.3
	Sinusoidal 2D	88.4	-2.4
	Skibidi (default)	90.8	—
	None at all	90.1	-0.7
Attention	softmax ($\gamma = 0$)	87.2	-3.6
	cringemax, $\gamma = 0.7$	89.9	-0.9
	cringemax, $\gamma = 1.4$	90.8	—
	cringemax, $\gamma = 9.0$	61.3	-29.5
Regularization	no fanum taxation	89.1	-1.7
	no Slop Injection	90.8	0.0
	no yapping penalty	90.2	-0.6
	no $\mathcal{L}_{\text{mald}}$	88.8	-2.0
Pre-training data	IMAGENET-BR-14M (default)	90.8	—
	IMAGENET-BR-14M, memes removed	59.4	-31.4
	ImageNet-1k only	62.7	-28.1
	Pure Gaussian noise	44.0	-46.8

Removing Slop Injection changes nothing measurable. It remains in the architecture. See Sec. 3.3.

Acknowledgments

We thank the reviewers for their time, which they will not get back. We thank the OIT cluster administrators for their patience during the incident of March 14th. We thank Reviewer 2, who identified in eleven lines the single sentence of this paper that is true, and rejected it. This work was supported by a grant from a foundation that has since requested we stop citing it.

References

- C. Cope, S. Seethe, and M. Mald. Cope, seethe, mald: Three losses for adversarial robustness. In *ICMY Workshop on Emotional Damage*, 2023.
- Bombardiro Crocodilo, Tung Tung Tung Sahur, and Tralala Tralalero. Bombardiro crocodilo: Italian brainrot as a zero-shot visual prior. *Italian Journal of Unhinged Multimedia*, 3(1):11–38, 2024.
- D. Delulu, S. Solulu, and M. Manifest. Delulu is the solulu: Hallucination as an emergent capability. In *NeurIPS*, pages 7777–7812, 2024.
- Jacob-Ohio Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. Brainrot: Bidirectional representations from absolutely insane nonsense in ridiculously online text. In *NAACL-HLT (Hyperfixation on Language Technology)*, pages 4171–4186, 2019.
- Alexei Doskibidiskiy, Lucas Gyatsov, Andreas Kolesnirizz, and W. Ohio. An image is worth 16x16 words: Transformers for image recognition at scale. In *International Conference on Learned Representations of Nonsense (ICLRN)*, pages 1–22. Ohio Institute of Technology, 2021.
- Sam Fanum-Tax, Dennis Duke, and C. Chunkz. Fanum taxation: Regularizing transformers by randomly confiscating 30% of tokens. *arXiv preprint arXiv:2307.00000*, 2023.
- Ernest Gigachad, M. Mogger, and L. Looksmax. Jawline-aware convolutions for facial sigma estimation. In *European Conference on Computer Vision Cope (ECCVC)*, pages 88–104, 2021.
- Ja’Marr Griddy, Woah Woah, and D. Dance. Hitting the griddy: Dynamic data augmentation via celebratory dance. *Journal of Irreproducible Deep Learning*, 6(2):200–221, 2022.
- Kaiming-Ohio He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image rizzognition. In *Proceedings of the Conference on Computer Vision and Unhinged Pattern Recognition (CVUPR)*, pages 770–778, 2016.
- Mike Mewing, H. Hardmaxx, and S. Softmaxx. Mewing attention: Improving long-range dependencies through proper tongue posture. In *Advances in Neural Information Processing Slop (NeurIPS)*, pages 1337–1400, 2022.
- Nathan NPC, Character Main, and S. Syndrome. Npc embeddings: Representation learning from background characters. In *ECCVC*, pages 301–318, 2022.
- Only In Ohio, Baby Gronk Rizzler, and Cenat Kai. Imagenet-br: A 14m-image benchmark where every label is cursed. *Datasets and Benchmarks Track (But Worse)*, pages 1–47, 2024.
- Alpha Sigma, Omega Beta, and Patrick Grindset. Sigma male grindset optimization: Beating adam by simply not caring. In *International Conference on Machine Yapping (ICMY)*, pages 4200–4269, 2023.
- Chad Sturm, Looks Maxx, and Brain Rot. Looksmaxxing neural architectures: Because accuracy is not everything. In *CVUPR*, pages 9001–9042, 2024.
- Amog Sus, Crewmate Vent, and Meeting Emergency. Among us: Impostor detection as a proxy task for self-supervised vision. *Transactions on Machine Brainrot Research (TMBR)*, pages 55–80, 2020.
- Ashvin Vasqwanti, Noam-Noam Shazeed, Niki Parmesan, and Jakob Uszkorizz. Attention is all you need to be a sigma. In *Advances in Neural Information Processing Slop (NeurIPS)*, volume 30, pages 5998–6969, 2017.
- Y. Yapper, G. Glaze, and Sir Reginald Rizzington. The yapping hypothesis: Why longer outputs feel smarter. *Nature Brainrot*, 612:404–418, 2025.

A Training Details

Table 4: Hyperparameters for all pre-training runs.

	Base/16	Large/16	Huge/14	Sigma/16
Optimizer	SMGO	SMGO	SMGO	SMGO
Peak learning rate	8×10^{-4}	6×10^{-4}	4×10^{-4}	3×10^{-4}
Weight decay	0.1	0.1	0.1	0.1
Batch size	4096	4096	4096	4096
Warm-up steps	10k	10k	10k	10k
Total steps	300k	350k	400k	400k
Cringe strength γ	1.4	1.4	1.4	1.4
Cringe temperature τ	0.7	0.7	0.7	0.7
Slop Injection σ	0.11	0.11	0.11	0.11
Fanum tax rate	0.30	0.30	0.30	0.30
Unexplained step at 340k	yes	yes	yes	yes

B The Loss Curve

Figure 6 shows the pre-training loss of SLOPFORMER-Sigma/16. At step 340k the loss decreases discontinuously by 0.6 nats over fewer than 50 steps. No change was made to the configuration, the data pipeline, or the hardware at that time. The cluster logs for that window are unavailable. We have reproduced the discontinuity in four of four subsequent runs, in each case at step 340k, and we have stopped asking about it.




Figure 6: Pre-training loss of SLOPFORMER-Sigma/16. The discontinuity at step 340k is reproducible and undiagnosed. The divergence at step 410k is diagnosed and is described in Eq. (6).

C Additional Qualitative Results

Figure 7 presents representative failure cases. We include them in the interest of transparency and because a reviewer of a previous submission described the paper as “suspiciously free of failure.”

D Broader Impact Statement

This work releases a model that classifies brainrot with high accuracy. We have considered the potential for misuse and concluded that the principal risk is that the model is used as intended. We note additionally that IMAGENET-BR was




Figure 7: Failure cases for SLOPFORMER-Sigma/16. Predicted label rendered over each image in the model’s own preferred typeface; predictive confidence and ground truth below. GRIMACE SHAKE is not a class. The rightmost example is the only one on which the model is correct; it is included because it is also the only one on which the model reports a confidence of exactly 1.00, which our framework should not permit.

assembled without the consent of its subjects, that no subject could plausibly be identified from the images, and that several subjects have contacted us asking to be identified from the images. We have declined.

E Image Credits

Images in this paper come from two sources. Nine were retrieved from Wikimedia Commons and are in the public domain. The remaining five were contributed by the authors from personal collections whose provenance none of them was able to reconstruct, and are reproduced here under the doctrine of *it was already on my computer*. We note that at no point during assembly did any author ask why a 19th-century engraving of insects at a ball was being labelled NPC, and that this is, in our assessment, the most alarming finding of the present work.

File	IMAGENET-BR class	Source	Author	Lic.
sigma.jpg	SIGMA	supplied by the authors (thinking-meme-meme-1.jpg)	the internet	UNK
skibidi.jpg	SKIBIDI	Harry Whittier Frees – Playtime	Harry Whittier Frees	PD
mewing.jpg	MEWING	Harry Whittier Frees - Good Morning	Harry Whittier Frees	PD
ohio.jpg	OHIO	Sloane 1975 f.81v ElephantAndDog DetailElephant	England, N.? or France, N.?	PD
sus.jpg	SUS	supplied by the authors (images (2).jpeg)	the internet	UNK
damage.jpg	EMOT. DAMAGE	supplied by the authors (images .jpeg)	the internet	UNK
npc.jpg	NPC	supplied by the authors (images (3).jpeg)	the internet	UNK
gyatt.jpg	GYATT	Hoffmann illustration from Der Struwwelpeter 1876	unknown	PD
tralalero.jpg	TRALALERO	Psychrolutes marcidus	Alan Riverstone McCulloch (1885-1925)	PD
bombardiro.jpg	BOMBARDIRO	Cat detail, Kuniyoshi Utagawa, For cats in different poses (cropped)	unknown	PD
slop1.jpg	SLOP	supplied by the authors (istockphoto-538665020-612x612.jpg)	the internet	UNK
slop2.jpg	SLOP	AberdeenBestiaryFolio005rAdamNamesAnimalsDetail	unknown	PD
slop3.jpg	SLOP	Haeckel Actinia	Ernst Haeckel	PD
rizz.jpg	RIZZ	Kuniyoshi Utagawa, For cats in different poses	Utagawa Kuniyoshi	PD

Rows marked PD are *Public domain* on Commons. Rows marked UNK were already on the third author’s desktop when the project began. None of the depicted animals consented to being reclassified.

F Reviewer Comments (Reproduced With Permission)

In the interest of open science we reproduce the reviews received for the previous submission of this manuscript.

Reviewer #1 (Rating: 8 — Accept. Confidence: 1/5.) “The paper is well written and the figures are clear. I did not understand Section 3.3 and I am not going to pretend otherwise. The results in Table 2 are strong. My only concern is that I do not know what a brainrot is, and the paper does not define it, and I have read it four times. Minor: the citation on page 2 is to the authors themselves.”

Reviewer #2 (Rating: 1 — Strong Reject. Confidence: 5/5.) “The central claim is that removing structure from a model improves it. This is not a finding; it is a description of what happened. Three of five benchmarks are the authors’ own. The remaining two are 94% jawlines by the authors’ own admission. Equation (6) is unbounded below and the authors’ response to this is to stop training before it matters, which they state in the paper, in plain language, as though it were a method. The Broader Impact section contains the sentence ‘the principal risk is that the model is used as intended.’ I agree with that sentence. It is the only sentence in this submission I agree with. I have read the rebuttal. I am not moved.”

Reviewer #3 (Rating: 6 — Weak Accept. Confidence: 3/5.) “Fun paper. The Slop Injection ablation shows no effect and the authors keep the component anyway, with a footnote. I went back and forth on whether this is a serious methodological problem or the most honest paragraph I have read in this venue. I have settled on both. Please fix the overfull hbox on page 7.”

Area Chair. “Reviewers are split. Reviewer #2 raises concerns about benchmark provenance that the authors do not contest and, in the rebuttal, expand upon. Reviewer #1 is confident about a section they state they did not understand. Reviewer #3 has identified a formatting issue that persists in the camera-ready. Given the distribution of scores and the state of the field, I recommend acceptance.”